Assignment III: CUDA Basics II

Aritra Bhakat

November 29, 2023

GitHub repo

Exercise 1: Histogram and Atomics

1. **Describe all optimizations you tried regardless of whether you committed to them or abandoned them and whether they improved or hurt performance.**

   The first optimisation I tried (after implementing a global memory version), was to keep a local histogram in shared memory. Then, I performed a reduction between the blocks, adding the local bins atomically to the global histogram. I let each thread within a block be responsible for adding a region of the bins, since the max thread block size (1024) is smaller than the number of bins (4096). This did not work unfortunately and increased the time initially. Using 1024 threads gave the best performance for the method, but the simple global memory method outperformed by close to a factor of 2.

   To reduce contention in global memory I didn't perform the atomic add if a bin is empty. Even though it leads to branching it reduced the histogram kernel execution time.

   Given that the shared histogram kernel was not performing better than the global one, I wanted to test whether it could be due to the GPU architecture. Therefore I tested it on both my machine with a GTX 1070 and in Google Colab with the Tesal T4, using an input length of `2^28 = 268435456`. I measured the time taken by the histogram kernel in ms:

	GTX 1070	Tesla T4
Global	12.052	32.551
Shared	21.110	45.390

The shared version still performs worse than the global one, although its relative performance to the global implementation is better on the Tesla T4. I think the reason the shared implementation performs poorly, or rather that the global version performs well, is that the bin count is so large that contention is minimal. Therefore there is more overhead when creating a shared memory version of the program.

I then simplified the way I fill the bins in the threads (threads need to fill up multiple bins since blockDim.x < NUM_BINS) from changing a contiguous portion as such:

```
int bin_region_size = num_bins / blockDim.x;
int bin_fill_start = bin_region_size * threadIdx.x;
for (int i = bin_fill_start; i < bin_fill_start + bin_region_size; i++) {
    ...
}
```

to filling bin elements at an interval of `stride`:

```
int stride = blockDim.x;
int i = threadIdx.x;
while (i < num_bins) {
  ...
  i += stride;
}
```

Now, on the Tesla T4 the shared memory version actually performs better than the global one! It doesn't do so on the GTX 1070 though, but does perform slightly better than before. Here are the results in ms:

	GTX 1070	Tesla T4
Shared (improved)	18.567	22.986

I'm not sure why the change doubled the speed for the Tesla T4, while there only is a minor improvement for the GTX 1070. It could be that it reduces the amount of cache misses, since the warps within a block are in sync and are accessing data from bins close to each other in shared memory.

2. **Which optimizations you chose in the end and why?**

   I stuck with the shared memory implementation and all the optimisations mentioned previously for the purpose of the assignment, even though it performs worse than the global memory one for GTX 1070.

3. **How many global memory reads are being performed by your kernel? Explain**

- Each thread reads its corresponding input data element from global memory
  - `num_elements` reads
- Each block updates all the bins in global memory
  - up to `num_bins * gridSize` reads (a check is performed to reduce unnecessary atomic operations to reduce contention)

4. **How many atomic operations are being performed by your kernel? Explain**

   Up to `num_bins * gridSize` atomic operations are performed by the kernel, as each block adds it's own local histogram to the global one. Atomic operations are avoided when bins are empty to reduce contention.

5. **How much shared memory is used in your code? Explain**

   Each block uses `sizeof(unsigned int) * NUM_BINS = 16.38 kB` of shared memory to store the shared histogram.

6. **How would the value distribution of the input array affect the contention among threads? For instance, what contentions would you expect if every element in the array has the same value?**

   Having many of the same value in the input array will increase contention between threads, since they will need to wait for more threads performing atomic addition to the same bin in global memory. If every element in the array has the same value, the kernel would degenerate into serial code.

7. **Plot a histogram generated by your code and specify your input length, thread block and grid.**

   The histogram was plotted for an input length of $n = 2^{15}$, using `1024` threads per block, and `(num_elements + TPB_HIST - 1) / TPB_HIST = 32` blocks.

8. **For a input array of 1024 elements, profile with Nvidia Nsight and report Shared Memory Configuration Size and Achieved Occupancy. Did Nvsight report any potential performance issues?**

```
Shared Memory Configuration Size    Kbyte   32.77
...                                 ...     ...
Achieved Occupancy                  %       47.57
```

   No potential issues were reported by Nsight.

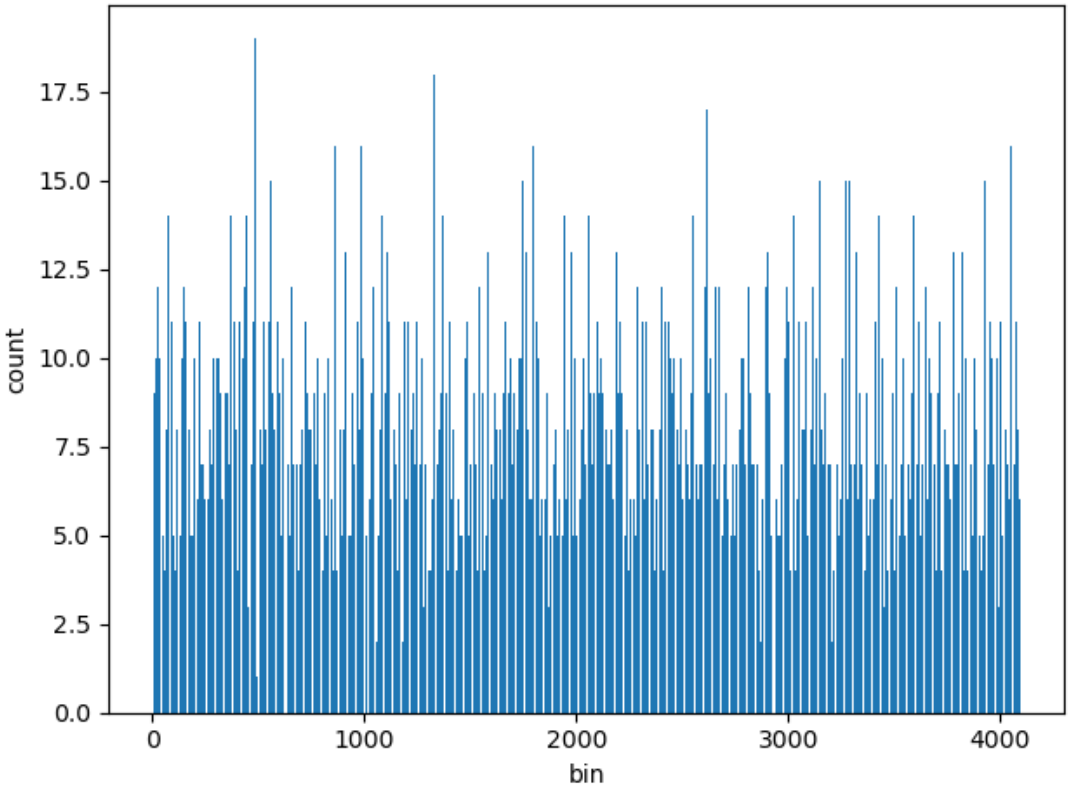

Figure 1: Histogram

Exercise 2: A Particle Simulation Application

1. **Describe the environment you used, what changes you made to the Makefile, and how you ran the simulation.**

   I use a WSL2 Ubuntu environment within my Windows PC. I had to change the Makefile variable for the GPU architecture to `ARCH=sm_61` which corresponds to my GTX 1070 GPU architecture. I also created two variables in the Makefile, `DEVICE` and `DEBUG`, to make it easier to compile on either CPU or GPU and also in debug mode so I could use `cuda-gdb`. I could then run the executable with my chosen configuration. I also had to modify the build structure so that I could for instance quickly switch between the CPU and GPU build.

2. **Describe your design of the GPU implementation of mover_PC() briefly.**

   I created a kernel for both `mover_PC` and `interP2G` for my port. Rather than doing all the allocation within the GPU versions before calling the kernel, I ported the allocation code for `particle`, `EMfield`, `grid`, `interpDensSpecies` and `parameters` to CUDA. That includes functions for performing synchronisation between the host and device.

   The `part` structure only needs to be synced once, at the start of the program. This is because the particles are never directly accessed by the host code. Instead, when we are in the `interP2G` kernel we access the `part` struct on the device, and use that to update the `ids` values. This does need to be synced back to the host, as there are some host side function like `applyBCids` accessing. It also modifies the data, so it needs to be synced to device at the start of every cycle. The, `field`, `grid` and `param` structs are all constant, so only need to be synced once at the start.

   For allocation, I needed to not only copy the structs, but perform a deep copy of all the data arrays that are pointed to in the flat arrays, and make sure the pointers on the device version of the struct point to the correct data.

   In order to retain the chained pointer functionality, I also ported the relevant parts of `Alloc.h` to CUDA. That involved allocating multiple levels of struct that hold pointers to pointers to pointers etc. And then correctly calculating the pointer offset.

   This allowed me to mostly leave the kernel code unchanged from the original host version of `mover_PC` and `interP2G`, as I didn't have to change any indexing. The main change was in the `interP2G` kernel where I had to use `atomicAdd` when accessing data in `ids`, to avoid race conditions from concurrent access.

3. **Compare the output of both CPU and GPU implementation to guarantee that your GPU implementations produce correct answers.**

   I used ParaView to visually compare the two implementations visually, and qualitatively there seemed to be no difference for both the `GEM_2D` and `GEM_3D` inputs.

I also compared the output values between the versions, to find the maximum difference between the results from the CPU and GPU. For `GEM_2D` we get:

```
B_10.vtk: loaded 98301 values
Max discrepancy: 0.0

E_10.vtk: loaded 98301 values
Max discrepancy: 0.0

rhoe_10.vtk: loaded 32768 values
Max discrepancy: 1.0000000000287557e-07

rhoi_10.vtk: loaded 32768 values
Max discrepancy: 1.0000000000287557e-07

rho_net_10.vtk: loaded 32768 values
Max discrepancy: 1.040000000000112e-07

Overall max discrepancy: 1.040000000000112e-07
```

We can see the difference is negligible, and comes down to floating point errors. There is no difference in the E and B file since they are constant and do not change after initialisation. Similarly for GEM_3D we get:

```
...

rho_net_6.vtk: loaded 524288 values
Max discrepancy: 1.0000000000287557e-07

E_10.vtk: loaded 1572861 values
Max discrepancy: 0.0

E_6.vtk: loaded 1572861 values
Max discrepancy: 0.0

B_9.vtk: loaded 1572861 values
Max discrepancy: 0.0

Overall max discrepancy: 1.0999999999934673e-07
```

4. **Compare the execution time of your GPU implementation with its CPU version.**

   Using the built in execution timer in the program we get the following execution time in seconds:

	GEM_2D	GEM_3D
CPU	48.3993	232.455
GPU	4.69184	65.8872

However, writing the result to disk is included in the timing, which feels a bit unfair. Especially given that the `GEM_3D` input saves large files to disk in every cycle. Therefore I modified the inputs and removed writing to disk for the purpose of performance measurements, and giving a more fair comparison of the running time between the implementations.

	GEM_2D	GEM_3D
CPU	46.9711	186.373
GPU	4.41544	14.776

For `GEM_2D` we obtain a 10.6 times speedup, and `GEM_3D` 12.6 times speedup.